

# Exploring Zero-Shot Learning Capabilities of LLMs

Gaurav Gulati

gauravgulati@gmail.com

Dronacharya College of Engineering  
India

Prof. Parveen Kumari

Dronacharya College of Engineering  
India

Garima Rana

gsr15597@gmail.com

Dronacharya College of Engineering  
India

Dr. Ritu Pahwa

Dronacharya College of Engineering  
India

## ABSTRACT

Transformative tools in natural language processing have recently emerged in the form of large language models (LLM), which have resulted in substantial advances in zero-shot learning (ZSL)[11]. Large language models can apply their broad knowledge base to accomplish diverse tasks that they have not been trained on, all without requiring specialized data for each labeled task.[24] ZSL research is a crucial area of machine learning because of its potential to tackle challenges in fields where annotated datasets are limited. The performance of zero-shot learning is significantly affected by the architecture of the models, the variety of pre-training datasets used, and the natural language prompt design.

In this study, systematic experimentation was used to investigate these factors using established benchmark datasets. The proposed methods assess the capabilities of LLMs in natural language inference, sentiment classification, and summarization by using specially designed prompts to optimize performance. This research examined the influence of model size and proposed techniques to mitigate the prevalence of spurious results commonly observed in zero-shot contexts.

The study indicates that larger models and diverse training data substantially improve zero-shot learning results, with well-designed prompts[16] being crucial for maximizing the capabilities of large language models. In addition, we identify key trade-offs that could result in unforeseen outcomes and propose remedies to address these challenges.

This study provides substantial information on the capabilities and limitations of big language models in zero-shot learning[6]. This framework facilitates the effective and responsible use of LLMs by outlining their benefits and drawbacks, enabling their application in solving real-world problems in the absence of labeled data.

## KEYWORDS

Large Language Models (LLMs), Zero-Shot Learning, Few-Shot Learning, Natural Language Processing (NLP), Prompt Engineering

## 1 INTRODUCTION

The problem of propagating models from little data continues to plague the areas of machine learning and artificial intelligence. One novel strategy is Zero-Shot Learning (ZSL), which enables models to complete tasks without the need for explicit training examples. The rise of Large Language Models (LLMs)[24] has sped up development in this area by offering previously unheard-of chances to finish challenging tasks with less oversight. The dynamic capabilities of large language models in zero-shot learning are examined in

this research, which focuses on the substantial knowledge gap in artificial intelligence by examining the models' performance across a variety of unexpected tasks.

Comprehending the potential and limitations of Zero-Shot capabilities is crucial for the development of adaptable and scalable AI systems. Through this study, a clearer knowledge of how prompt engineering and LLM[28] [20] design affect an LLM's endurance to tackle formerly unlabeled jobs will be acquired. The research conclusions have important implications with wide-ranging real-world implications in machine learning, natural language processing, and AI ethics.

Evidently, large language models (LLMs) exhibited impressive accomplishments in sentiment analysis, summarization, and natural language inference[25], with some models achieving these tasks with little to no specialized fine-tuning. Addressing task-specific nuances and ensuring consistent performance across many regions are only two of the ongoing problems. To address these issues, this study meticulously examined various model topologies, pre-training datasets, and prompt design techniques to improve zero-shot learning performance.

To address these discrepancies, extensive experiments are conducted on multiple standardized test datasets to evaluate the performance of different Large Language Models on multiple tasks, and investigate the impact of different pretraining data and prompting strategies.

The linkage between model design, dataset variability, and prompt techniques is examined in this paper, with a focus on how these factors affect the performance of Zero-Shot Learning.

Not only this discussion advance our knowledge of LLMs' zero-shot learning capabilities, but it also has the potential to impact the construction of AI applications. The intended goal of the present study is to enhance the efficacy of models in practical environments when task-specific data is likely nonexistent or insufficient.

**In Zero-Shot Learning (ZSL) [4], machine learning models are trained without labeled data to perform tasks that have never been attempted before.** The foundation of ZSL is the utilisation of previously acquired knowledge, which is then expanded to handle new tasks using contextual or semantic information provided during the inference stage, usually in the form of instructions in natural language.[19]

**Few-Shot Learning (FSL)[7] addresses scenarios where a novel job has a restricted number of annotated examples available.** These sparse examples serve as the foundation for FSL's model, which extrapolates the task's fundamental structure and expands its application to new situations. The input prompt usually

contains examples, that allow the model to recognise patterns with comparatively little guidance.

Both ZSL and FSL[10] leverage the power of pre-trained Large Language Models, which allows them to function well in situations when resources are scarce and labeled data is difficult to obtain. While FSL offers a fundamental framework that gives guidance, ultimately improving comprehension and task performance, ZSL demonstrates the model's capacity to generalise beyond particular examples. These theoretical foundations are responsible for the capacity for generalisation of artificially intelligent systems in a wide range of fields, and they are especially useful when gathering labeled data is difficult or costly.

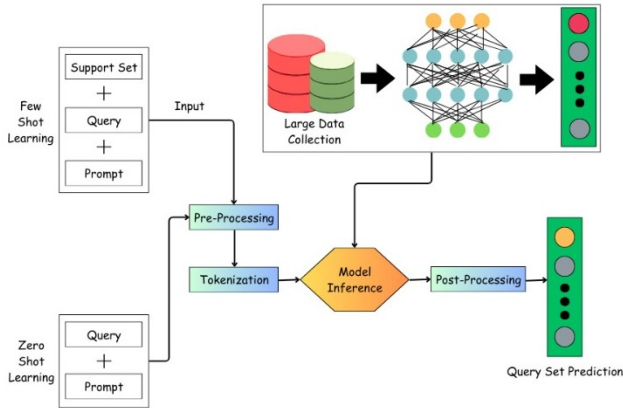


Figure 1: Zero-shot Inferencing Pipeline

## 2 EXPERIMENTAL SETUP

For this study, the primary computing environment was the **Kaggle**[13] platform. With its comprehensive cloud-based setup, Kaggle removes the need for local hardware restrictions while serving the resources indispensable for processing and developing models. The platform's smooth communication with several libraries and GPU availability greatly simplified the model's evaluation process.

The models were developed and validated using Kaggle Notebooks, which greatly accelerated the workflow and facilitated remote code execution. Throughout the experimentation phase, optimal performance was guaranteed by the lack of hardware restrictions and the capacity to scale up tools on Kaggle.

The **Hugging Face Transformers**[8] library was integrated, offering an extensive assortment of pretrained models, such as **Qwen2.5-1.5B-Instruct**[22], **Llama-3.2-3B-Instruct**[29], and **Gemma-2-2B-it**[27], for model construction and assessment. These models, which emphasized inferences regardless of prior task-specific training, were optimized for zero-shot learning tasks. For the sake of efficient computation and a seamless relationship with Hugging Face's APIs, **PyTorch**[5] was utilized for model setup and management.

By utilizing Kaggle's resources, the experiment was carried out effectively and without regard to hardware constraints, offering a strong basis for investigating the potential of huge language models in natural language processing.

## 3 LIBRARIES

To evaluate the zero-shot learning capabilities of large language models, this study employed essential frameworks, including PyTorch and Hugging Face Transformers. For natural language processing, the Hugging Face Transformers library offers pre-trained models and tools, while the PyTorch library offers an adaptable and effective framework for putting these language models into practice and executing.

### 3.1 Hugging Face

For an expansive variety of natural language processing (NLP)[17] applications, the Hugging Face Transformers[8] library showcases an exceptionally strong open-source framework. Numerous pre-trained models and tools have been made accessible through this library, assisting tasks like text generation, question answering, and zero-shot classification[2]. With the help of its many capabilities, we could quickly implement cutting-edge language models without needing a lot of computational or training time.

Hugging Face makes it easier to integrate and use sophisticated models by offering pre-built pipelines and tokenizers that enable inference operations to be executed directly. Through its smooth integration with frameworks such as PyTorch, computational resources are efficiently controlled, allowing for the study of large language models in a variety of natural language processing scenarios. Due to its versatility and ease of use, Hugging Face proved to be an indispensable resource for the research on zero-shot learning.

### 3.2 PyTorch

A further significant component of this study was PyTorch[5], an acclaimed open-source machine learning system. PyTorch, a framework esteemed for its flexibility and efficiency in deep learning applications, supplies robust features for dynamic computational graphs, automatic differentiation, and tensor computation. The aforementioned features enabled the precise and efficient construction and evaluation of large language models.

High-performance inference tasks in this study were facilitated through the use of PyTorch for loading and executing the selected pre-trained models. Its interoperability with Hugging Face Transformers made it possible for the libraries to integrate seamlessly, guaranteeing that zero-shot learning[9] procedures performed smoothly. PyTorch's dynamic computation graph feature proved especially useful for managing intricate models and maximizing computational effectiveness when conducting experiments.

A simplified process for putting large language models into practice and assessing them was made possible by the merging of Hugging Face Transformers and PyTorch. In aggregate, these libraries provide the computational efficiency, scalability, and flexibility needed to thoroughly investigate models' zero-shot learning potential. This combination became a crucial component of the research framework, as it allowed for the effective handling of the challenges posed by various NLP[17] tasks.

## 4 DATA COLLECTION

Assessing the large language models' (LLMs') zero-shot learning [9] capabilities required careful consideration of the data collecting procedure. This section describes the data collection, processing, and

organization approach, which includes using the **Google Boolean Question (BoolQ) Dataset**[3] in a structured format consisting of a question and a Boolean response for efficient model evaluation.

#### 4.1 Data Capture Setup

The **BoolQ** dataset, which consists of yes/no questions combined with contextual excerpts, served as the main dataset for this investigation. In order to make sure it was accurate and in accordance with the input criteria of the designated models, this dataset was pre-processed. This configuration offered a solid framework to rank LLMs' classification and reasoning abilities in a zero-shot scenario[21], accentuating their knack to perceive and adjust to contextual information.

#### 4.2 Data Format and Sample Collection

BoolQ is a question-answering dataset designed for yes/no questions, comprising 15,942 examples, where each sample is a triplet of (question, passage, answer); the questions are naturally occurring, generated in unprompted and unconstrained contexts, with a binary label of "FALSE" or "TRUE" as the response.

For illustration: **"Is the sky blue?"**, **"The sky appears blue when viewed from Earth during daylight hours."**, **"TRUE"**

In accordance to this format, the inquiry appears first, then the context with pertinent details, and finally the binary label with the response. A wide range of questions and instances are incorporated in each example so that we had an adequate base of data to judge our algorithms' zero-shot learning abilities.

#### 4.3 Challenges and Considerations

During the evaluation of our zero-shot learning models[6], we encountered several problems that required careful consideration. Assuring consistent performance across many datasets was one of the main challenges; In particular, the models did not achieve the expected accuracy on the GSM8K[15] dataset. The intricacy of applying larger models to specific tasks without fine-tuning was brought to light by this constraint, which occasionally produces less-than-ideal results for specialized or smaller datasets.

Considering these constraints, smaller models like **"qwen2.5-1.5b"**, **"llama3.2-3b"**, and **"gemma-2-2b-it"** were selected instead of larger models, which could have produced more precise results but would have needed a lot more computing power. Notably on the Kaggle platform, where computational assets are constrained[14], the smaller models rendered a better trade-off between performance and resource usage. The selection of these models was contingent on their propensity for successfully carrying out zero-shot learning tasks having reasonable amount of resources.

Verifying the models' resilience against various input queries presented a further challenge. Despite having pre-trained models, our models occasionally experienced trouble with more complex reasoning tasks since real-time inference without fine-tuning in zero-shot settings was required. This concentrated on the trade-off between efficiency and accuracy when working with large, pre-trained models in a zero-shot setting.

All things taken into consideration, the models offered a firm foundation for determining zero-shot learning capabilities. Nevertheless, the decision was made to opt for smaller models due to the

specialized nature of the datasets and limited platform resources. This choice allowed for efficient exploration of zero-shot learning while optimizing the available resources.

## 5 RESULTS

The BoolQ dataset was used to assess the zero-shot learning models' performance using the following important metrics: accuracy, precision, recall, and F1-score.[1] These metrics offer a thorough assessment of the models' accuracy in classifying the data while maintaining a balance between true positive and false positive rates. For a thorough comparison, the categorization reports for each model are described in Tables 1, 2, and 3 illustrating both their advantages and disadvantages.

### 5.1 Classification Performance[12]

**Table 1: Classification Report for Qwen**

Class	Precision	Recall	F1-Score	Support
FALSE	0.74	0.70	0.72	3553
TRUE	0.82	0.85	0.84	5874
<b>Accuracy</b>	0.79 (9427 samples)			
<b>Macro Avg</b>	0.52	0.52	0.52	9427
<b>Weighted Avg</b>	0.79	0.79	0.79	9427

**Table 2: Classification Report for Gemma**

Class	Precision	Recall	F1-Score	Support
FALSE	0.81	0.71	0.76	3553
TRUE	0.84	0.90	0.87	5874
<b>Accuracy</b>	0.83 (9427 samples)			
<b>Macro Avg</b>	0.82	0.81	0.81	9427
<b>Weighted Avg</b>	0.83	0.83	0.83	9427

**Table 3: Classification Report for Llama**

Class	Precision	Recall	F1-Score	Support
FALSE	0.51	0.98	0.67	3553
TRUE	0.97	0.44	0.60	5874
<b>Accuracy</b>	0.64 (9427 samples)			
<b>Macro Avg</b>	0.50	0.47	0.43	9427
<b>Weighted Avg</b>	0.80	0.64	0.63	9427

Different models showed different levels of performance. The Qwen model's[23] weighted F1-score was 0.79, and its accuracy was 79%. With a weighted F1-score of 0.83 and an accuracy of 83%, the Gemma model[26] fared better than Qwen. Despite its excellent precision for the TRUE class, Llama[18] presented a weighted F1-score of 0.63 and a lower overall accuracy of 64%.

These findings demonstrate the trade-offs between consistency and accuracy in various zero-shot learning algorithms. Gemma had the best overall performance, while Qwen yielded competitive scores, and Llama showed promise in managing particular situations, despite recall issues for several classes.

## 5.2 Model Architecture Differentiation

The architectures as well as applications of the three models employed for this project fluctuate extensively:

- **Qwen2.5-1.5B-Instruct:**
  - **Compact Transformer Architecture:** This kind of transformer is ideal for instruction-following tasks and possesses a simplified architecture with 1.5 billion parameters.
  - **Multilingual Training and RLHF:** To strengthen its alignment with a range of human preferences, the model incorporates multilingual training capabilities and makes use of Reinforcement Learning from Human Feedback (RLHF).[30]
  - **Optimized for Efficacy:** Qwen2.5-1.5B-Instruct prioritizes low latency and computational economy, making it especially suitable for deployment in contexts with limited resources, assuring quick and economical performance.
- **Llama-3.2-3B-Instruct:**
  - **Dense Transformer Architecture:** With its robust 3 billion parameter architecture, this model provides impressive generalization capabilities for a number of tasks.
  - **Extensive Pretraining and Fine-Tuning Support:** It is adaptable for specific applications due to its extensive pretraining across a variety of domains, as well as its support for parameter-efficient fine-tuning techniques[31] like LoRA and rich natural language processing.
  - **Versatility with Higher Resource Demand:** Designed for a variety of NLP tasks, Llama-3.2-3B-Instruct balances capacity and resource considerations by delivering good performance at a higher computational cost than smaller models.
- **Gemma-2-2B-it:**
  - **Generalized Multi-Task Model Architecture:** This model has a flexible architecture with two billion parameters that is tailored for zero-shot and multi-task learning situations.
  - **Effective Inference and Multi-Task Pretraining:** Gemma-2-2B-it ensures uniform and efficient performance on a range of activities by utilizing multi-task pretraining in conjunction with excellent inference approaches.
  - **Strength in Classification and Reasoning:** The model is very flexible for dynamic, on-the-fly jobs and performs exceptionally well in reasoning and classification tasks, with a focus on real-time zero-shot applications.

Table 4: Comparative Summary of Model Architectures

Feature	Qwen	Llama	Gemma
Parameters	1.5B	3B	2B
Type	Compact	Dense	Multi-Task
Focus	Efficiency	Generalization	Reasoning
Strength	Low Latency	Versatility	Multi-Task
Limitations	Niche Tasks	Resource-Heavy	Complexity
Best Fit	Edge/Cloud	NLP Diversity	Zero-Shot Tasks

The trade-offs among the three models' efficiency, generalization, and task-specific performance are highlighted in this comparison.

On the Kaggle platform, smaller models like Qwen were selected due to their computational efficiency, whereas Gemma offered well-balanced multitasking capabilities. Llama's greater computational requirements limited its generalization potential. The models utilized in this project were all pre-trained for zero-shot inference tasks and include a total of 1.5 billion to 3 billion parameters. These models use revolutionary techniques for natural language processing that have been optimized for generalization across a range of NLP tasks.

## 5.3 Resource Utilisation

On the Kaggle platform, the models demonstrated efficient performance with reduced latency during inference, optimizing resource usage. Even though the models had a lot of parameters (1.5 billion to 3 billion), the zero-shot learning tasks were handled well, guaranteeing timely and precise answers. The models' capacity to manage intricate natural language tasks with little computing overhead is demonstrated here, which qualifies them for real-time inference in scalable cloud-based environments such as Kaggle.

## 6 CONCLUSION

This study highlights the intriguing potential of utilizing large language models for real-time zero-shot learning problems. Effective performance on the Kaggle platform was achieved by integrating the Hugging Face Transformers library and leveraging models such as Qwen2.5-1.5B-Instruct, Llama-3.2-3B-Instruct, and Gemma-2-2B-it. Without requiring task-specific fine-tuning, the models demonstrated their versatility in generalizing across a broad range of natural language processing tasks by doing exceptionally well on inference tasks.

The models produced precise and fast responses for zero-shot reasoning problems, yielding very excellent results. This demonstrates that even with challenging NLP problems, large pre-trained models can do high-performance inference in real-time.

There are a number of intriguing prospects to advance the models' capabilities in the future. To expand the models' applicability, future work may explore additional zero-shot learning challenges, enhance model efficiency through techniques such as quantization, or tailor the models for specific domains. Furthermore, the models' general performance and usefulness could be expanded by incorporating more complex multi-modal capabilities or optimizing them for specific use situations. This project paves the way for further research into large language models and zero-shot learning, spurring innovation in a number of fields.

## 7 ETHICAL CLEARANCE

We independently selected and processed the data for this research in order to assess the specified models' zero-shot learning capabilities. To ensure a thorough setup for testing the models' capacity to execute inference without fine-tuning, this method involved compiling a range of tasks, which were subsequently organized and labeled for use in our model evaluation.

By overseeing the labeling and data preparation phases, the process was ensured to be accurate and consistent. Additionally, the evaluation was conducted in a manner that preserved privacy, as the models used for inference did not rely on sensitive or personal

data. Moreover, employing pre-trained models from sources like Hugging Face guarantees that no sensitive user data is incorporated, as the models operate on general language tasks without the need for personally identifiable information.

## 8 ACKNOWLEDGMENT

We genuinely appreciate the essential suggestions as well as critiques from our mentor, Dr. Ritu Pahwa, and instructor, Prof. Parveen Kumari, which greatly improved the caliber of our work. We would want to give credit to the open-source community, particularly the creators of PyTorch, Hugging Face, and the pre-trained models that were utilized in this project. In order for us to deploy a successful zero-shot learning system for natural language processing jobs, their dedication to making machine learning technologies available and making resources easily accessible was crucial. Additionally, We also thank the authors of the research publications and sources that helped create and comprehend these potent models.

## REFERENCES

- [1] David M. W. Powers. 37–63. <https://doi.org/10.48550/arXiv.2010.16061>
- [2] Shashank Srivastava, Igor Labutov, and Tom Mitchell. 306–316. <https://doi.org/10.18653/v1/P18-1029>
- [3] Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. arXiv:1905.10044 [cs.CL] <https://arxiv.org/abs/1905.10044>
- [4] Edgar Schönfeld, Sayna Ebrahimi, Samarth Sinha, Trevor Darrell, and Zeynep Akata. arXiv:1812.01784 [cs.CV] <https://arxiv.org/abs/1812.01784>
- [5] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. arXiv:1912.01703 [cs.LG] <https://arxiv.org/abs/1912.01703>
- [6] Yongqin Xian, Christoph H. Lampert, Bernt Schiele, and Zeynep Akata. 2251–2265. <https://doi.org/10.1109/TPAMI.2018.2857768>
- [7] Yaqing Wang, Quanming Yao, James Kwok, and Lionel M. Ni. arXiv:1904.05046 [cs.LG] <https://arxiv.org/abs/1904.05046>
- [8] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander M. Rush. arXiv:1910.03771 [cs.CL] <https://arxiv.org/abs/1910.03771>
- [9] Omurhan A. Soysal and Mehmet Serdar Guzel. 1–4. <https://doi.org/10.1109/HORA49412.2020.9152859>
- [10] Suvarna Kadam and Vinay Vaidya. 100–112.
- [11] Yongqin Xian, Christoph H. Lampert, Bernt Schiele, and Zeynep Akata. arXiv:1707.00600 [cs.CV] <https://arxiv.org/abs/1707.00600>
- [12] Zeljko Vujovic. 599–606. <https://doi.org/10.14569/IJACSA.2021.0120670>
- [13] Teruaki Hayashi, Takumi Shimizu, and Yoshiaki Fukami. arXiv:2107.11929 [cs.CY] <https://arxiv.org/abs/2107.11929>
- [14] Xixi Li, Yun Bai, and Yanfei Kang. arXiv:2103.00501 [cs.SI] <https://arxiv.org/abs/2103.00501>
- [15] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. arXiv:2110.14168 [cs.LG] <https://arxiv.org/abs/2110.14168>
- [16] Victor Sanh, Albert Webson, Colin Raffel, Stephen H. Bach, Lintang Sutawika, Zaid Alyafeai, Antoine Chaffin, Arnaud Stiegler, Teven Le Scao, Arun Raja, Manan Dey, M Saiful Bari, Canwen Xu, Urmish Thakker, Shanya Sharma Sharma, Eliza Szczecula, Taewoon Kim, Gunjan Chhablani, Nihal Nayak, Debajyoti Datta, Jonathan Chang, Mike Tian-Jian Jiang, Han Wang, Matteo Manica, Sheng Shen, Zheng Xin Yong, Harshit Pandey, Rachel Bawden, Thomas Wang, Trishala Neeraj, Jos Rozen, Abheesht Sharma, Andrea Santilli, Thibault Fevry, Jason Alan Fries, Ryan Teehan, Tali Bers, Stella Biderman, Leo Gao, Thomas Wolf, and Alexander M. Rush. arXiv:2110.08207 [cs.LG] <https://arxiv.org/abs/2110.08207>
- [17] Diksha Khurana, Aditya Koli, Kiran Khatter, and Sukhdev Singh. 3713–3744. <https://doi.org/10.1007/s11042-022-13428-4>
- [18] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. arXiv:2302.13971 [cs.CL] <https://arxiv.org/abs/2302.13971>
- [19] Lei Wang, Wanyu Xu, Yihuai Lan, Zhiqiang Hu, Yunshi Lan, Roy Ka-Wei Lee, and Ee-Peng Lim. 2609–2634. <https://doi.org/10.18653/v1/2023.acl-long.147>
- [20] Yinheng Li. 641–647. <https://aclanthology.org/2023.ranlp-1.69>
- [21] Xuandong Zhao, Siqi Ouyang, Zhiguo Yu, Ming Wu, and Lei Li. 15590–15606. <https://doi.org/10.18653/v1/2023.acl-long.869>
- [22] Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. arXiv:2309.16609 [cs.CL] <https://arxiv.org/abs/2309.16609>
- [23] Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. arXiv:2308.12966 [cs.CV] <https://arxiv.org/abs/2308.12966>
- [24] Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, Rui Zheng, Xiaoran Fan, Xiao Wang, Limao Xiong, Yuhao Zhou, Weiran Wang, Changhao Jiang, Yicheng Zou, Xiangyang Liu, Zhangyue Yin, Shihan Dou, Rongxiang Weng, Wensen Cheng, Qi Zhang, Wenjuan Qin, Yongyan Zheng, Xipeng Qiu, Xuanjing Huang, and Tao Gui. arXiv:2309.07864 [cs.AI] <https://arxiv.org/abs/2309.07864>
- [25] Chen Ling, Xujiang Zhao, Jiaying Lu, Chengyuan Deng, Can Zheng, Junxiang Wang, Tanmoy Chowdhury, Yun Li, Hejie Cui, Xuchao Zhang, Tianjiao Zhao, Amit Panalkar, Dhagash Mehta, Stefano Pasquali, Wei Cheng, Haoyu Wang, Yanchi Liu, Zhengzhang Chen, Haifeng Chen, Chris White, Quanquan Gu, Jian Pei, Carl Yang, and Liang Zhao. arXiv:2305.18703 [cs.CL] <https://arxiv.org/abs/2305.18703>
- [26] Gemma Team. arXiv:2408.00118 [cs.CL] <https://arxiv.org/abs/2408.00118>
- [27] Gemma Team. arXiv:2403.08295 [cs.CL] <https://arxiv.org/abs/2403.08295>
- [28] Pittawat Taveekitworachai, Febri Abdullah, and Ruck Thawonmas. arXiv:2401.08273 [cs.CL] <https://arxiv.org/abs/2401.08273>
- [29] Yaowei Zheng, Richong Zhang, Junhao Zhang, Yanhan Ye, and Zheyang Luo. 400–410. <https://doi.org/10.18653/v1/2024.acl-demos.38>
- [30] Timo Kaufmann, Paul Weng, Viktor Bengs, and Eyke Hüllermeier. arXiv:2312.14925 [cs.LG] <https://arxiv.org/abs/2312.14925>
- [31] Venkatesh Balavadhani Parthasarathy, Ahtsham Zafar, Aafaq Khan, and Arsalan Shahid. arXiv:2408.13296 [cs.LG] <https://arxiv.org/abs/2408.13296>